

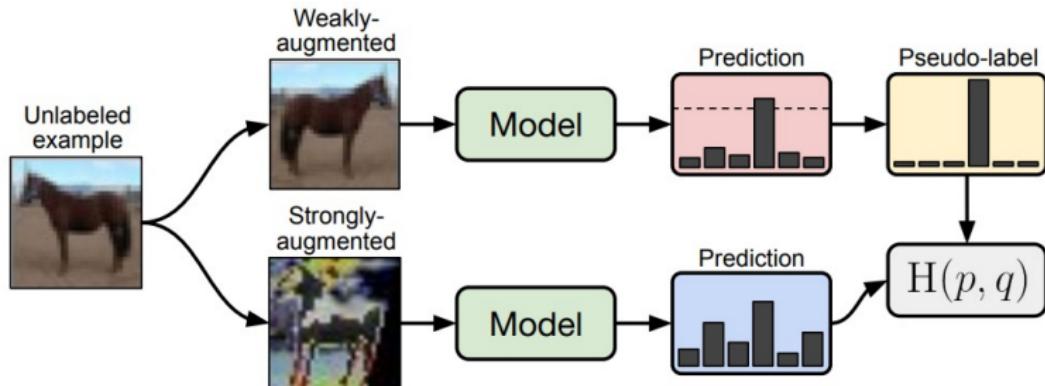
# FixMatch em CIFAR-10: Uma abordagem experimental

Tomás Lira, Gabriel Abad, Antônio Brych

FGV EMAp

13 de novembro de 2025

# Visão Geral do Modelo



**Figura:** Fonte: Sohn et al., 2025. FixMatch: Simplifying semi-supervised learning with consistency and confidence.

# Algoritmo do FixMatch

---

**Algorithm 1** FixMatch algorithm.

- 1: **Input:** Labeled batch  $\mathcal{X} = \{(x_b, p_b) : b \in (1, \dots, B)\}$ , unlabeled batch  $\mathcal{U} = \{u_b : b \in (1, \dots, \mu B)\}$ , confidence threshold  $\tau$ , unlabeled data ratio  $\mu$ , unlabeled loss weight  $\lambda_u$ .
  - 2:  $\ell_s = \frac{1}{B} \sum_{b=1}^B H(p_b, \alpha(x_b))$  {Cross-entropy loss for labeled data}
  - 3: **for**  $b = 1$  **to**  $\mu B$  **do**
  - 4:    $q_b = p_m(y | \alpha(u_b); \theta)$  {Compute prediction after applying weak data augmentation of  $u_b$ }
  - 5: **end for**
  - 6:  $\ell_u = \frac{1}{\mu B} \sum_{b=1}^{\mu B} \mathbb{1}\{\max(q_b) > \tau\} H(\arg \max(q_b), p_m(y | \mathcal{A}(u_b)))$  {Cross-entropy loss with pseudo-label and confidence for unlabeled data}
  - 7: **return**  $\ell_s + \lambda_u \ell_u$
- 

**Figura:** Descrição do algoritmo FixMatch (Sohn et al., 2020).

# FixMatch

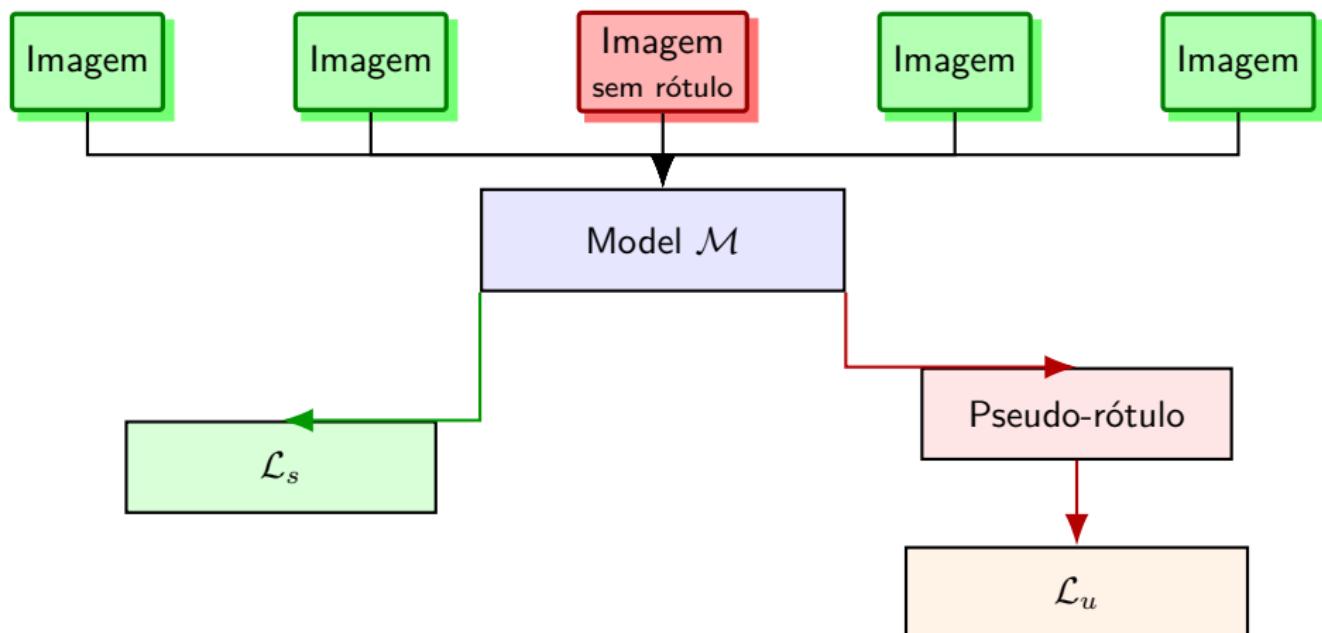
Intuição: Um bom modelo deve ser consistente, isto é, manter a mesma predição para uma mesma imagem sob augmentação forte e fraca.

O FixMatch se baseia nesse princípio para poder utilizar dados não rotulados em seu treinamento.

Questionamento: Por que calcular "pseudo-labels"em dados fracamente augmentados?

## Fluxograma (FixMatch): Exemplo em um 'batch'

4 verdes (rotuladas) e 1 vermelha (sem rótulo)



$\mathcal{L}_s$ : Perda supervisionada.

$\mathcal{L}_u$ : Perda não supervisionada.

# FixMatch

$$p_m(y|\mathbf{x}) = \text{softmax}(f_\theta(\mathbf{x}))$$

$$\hat{q}_b = \text{argmax}(p_m(y|\text{Augmentação Fraca}(\mathbf{x}_b)))$$

$$m_b = \max(p_m(y|\text{Augmentação Fraca}(\mathbf{x}_b)))$$

$$\mathcal{L}_s = \frac{1}{B} \sum^B H(y_b, p_m(y|\mathbf{x}_b))$$

Imposição de consistência

$$\mathcal{L}_u = \frac{1}{\mu B} \sum_{b=1}^{\mu B} \mathbf{1}(m_b \geq \tau) \underbrace{H(\hat{q}_b, p_m(y|\text{Augmentação Forte}(\mathbf{x}_b)))}_{\downarrow}$$

$$\mathcal{L}_{\text{total}} = \mathcal{L}_s + \lambda_u \mathcal{L}_u$$

# Estrutura Geral da Implementação

- A implementação foi totalmente guiada pelo paper original.
- Organização do código:
  - `make_cifar10_datasets()` – divide o CIFAR-10 em subconjuntos rotulados e não rotulados.
  - `LabeledWrapper` e `UnlabeledWrapper` – classes personalizadas que aplicam as transformações *weak* e *strong*.
  - `get_model()` – cria a **ResNet-18** ajustando a camada final para 10 classes.
  - `ModelEMA` – mantém a média exponencial dos pesos para avaliação estável.

## Função de Treino train\_fixmatch

- Cada **epoch** itera sobre batches rotulados e não rotulados simultaneamente.
- Estrutura principal:
  - ➊ **Forward supervisionado:** calcula a perda de entropia cruzada em  $(x_l, y_l)$ .
  - ➋ **Forward não supervisionado:** gera pseudo-rótulos a partir de  $(x_{ul}^{weak})$  e aplica a perda nas versões  $(x_{ul}^{strong})$ .
  - ➌ Filtra apenas as amostras com confiança  $\geq \tau = 0.95$ .
  - ➍ Soma ponderada:  $\mathcal{L} = \mathcal{L}_s + \lambda_u \mathcal{L}_u$  com  $\lambda_u = 1.0$ .
- Cada passo atualiza tanto o modelo principal quanto o modelo EMA.

# Detalhes das Aumentações

- **Aumentação Fraca (*weak\_transform*)**

- `RandomHorizontalFlip()` – espelhamento aleatório.
- `RandomCrop(32, padding=4)` – deslocamento até 12,5%.
- Normalização CIFAR-10: médias e desvios padrão padrão.

- **Aumentação Forte (*strong\_transform*)**

- `RandAugment(num_ops=2, magnitude=9)` – aumenta variação sem reduzir qualidade.
- `Cutout(length=8)` – remove regiões aleatórias da imagem (regularização espacial).

- A classe **ModelEMA** mantém uma cópia dos pesos atualizada a cada passo:

$$\theta_{\text{EMA}} \leftarrow \alpha \theta_{\text{EMA}} + (1 - \alpha) \theta \quad \text{com } \alpha = 0.999$$

- O modelo EMA é usado para avaliação final, garantindo previsões mais estáveis.
- EMA foi utilizado pois é uma recomendação, feita no paper original, para reportar os resultados.

# Decisões e Hiperparâmetros

- **Modelo:** ResNet-18 pré-definida do torchvision.
- **Augmentation:** RandAugment + Cutout (forte); Flip + Crop (fraco).
- **Limiar de confiança:**  $\tau = 0.95$
- **Otimizador:** SGD + Nesterov momentum (0.9)
- **Scheduler:** Decaimento cossenoide da taxa de aprendizado.
- **EMA:** Decaimento 0.999, usado apenas na avaliação.
- **Batch:**  $B = 64$ ,  $\mu = 7$

# Framework Experimental

Experimentos com 8 modelos distintos!

- **Família ResNet:** ResNet-18; ResNet-32; ResNet-34; Resnet-50; ResNet-101; ResNet-152;
- **Outros Modelos:** DenseNet-121; MobileNetV3

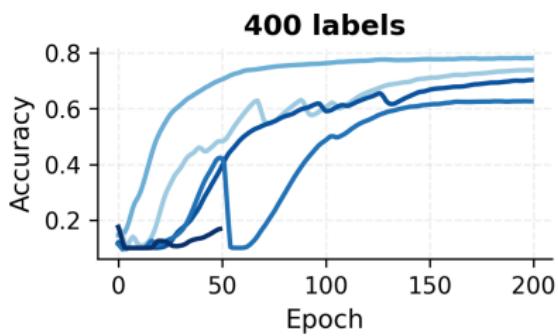
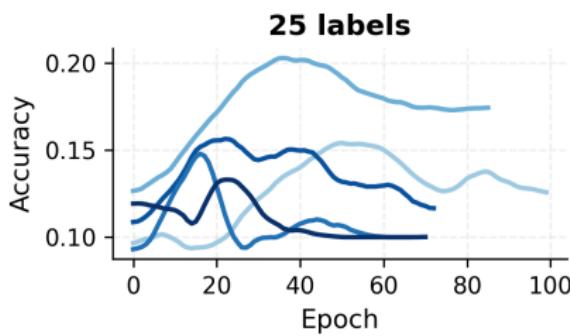
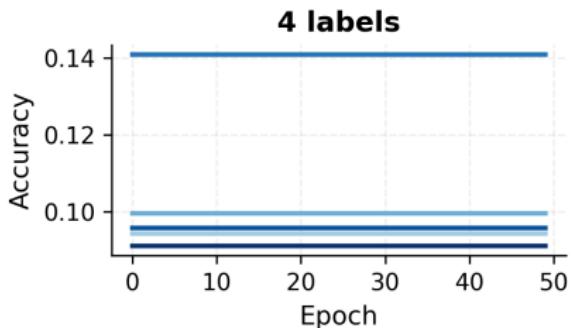
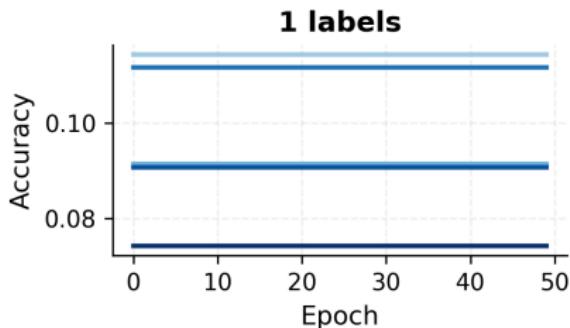
Testes realizados:

- Testes de Convergência
- Evoluções das curvas de acurácia
- Comportamento da 'Mask-Rate'

# ResNets

# Accuracy vs Epoch by Label Count

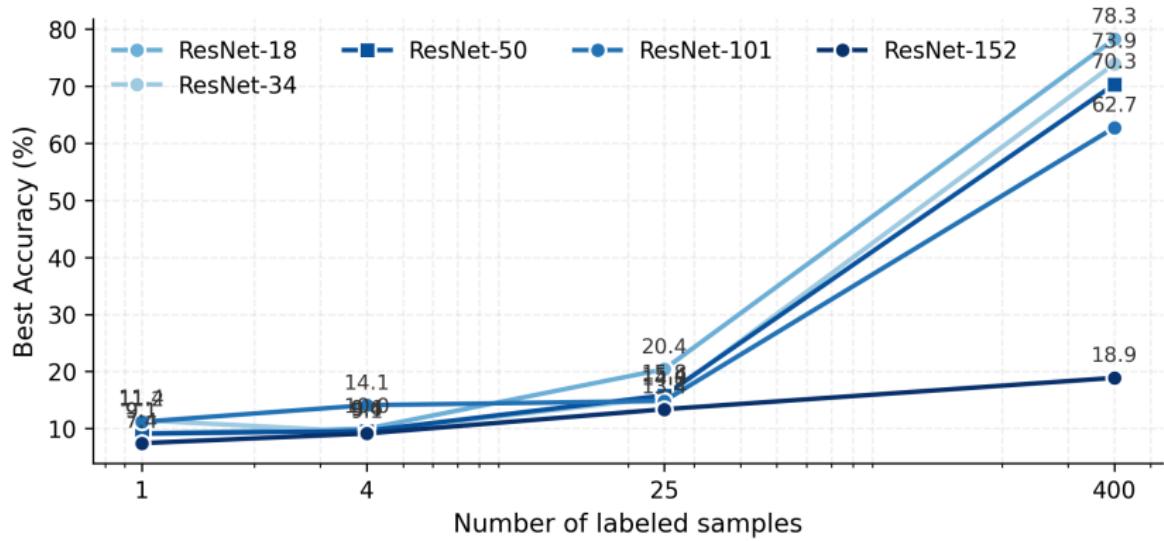
ResNet Family

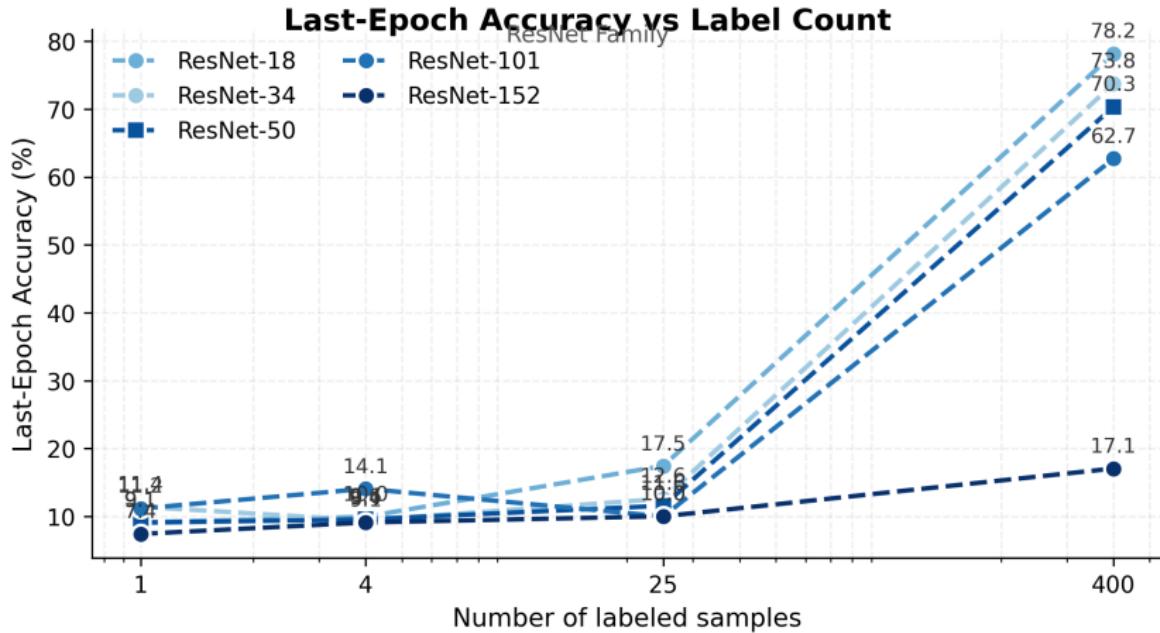


Legend:

- ResNet-18
- ResNet-50
- ResNet-101
- ResNet-152
- ResNet-34

## Best Accuracy vs Label Count



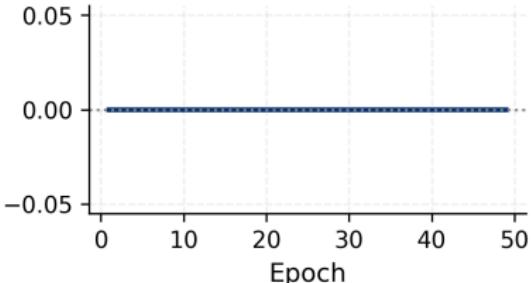


# Change in Supervised Loss vs Epoch by Label Count

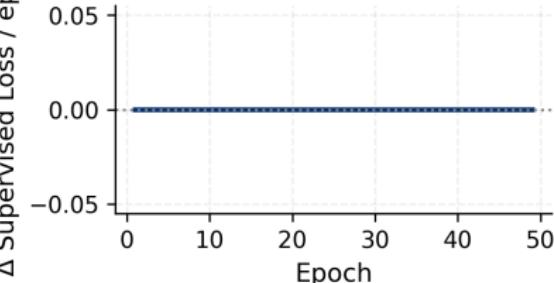
ResNet Family

$\Delta$  Supervised Loss / epoch

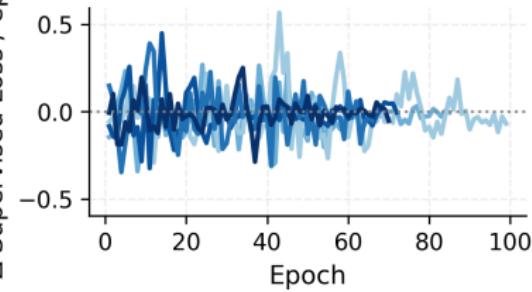
**1 labels**



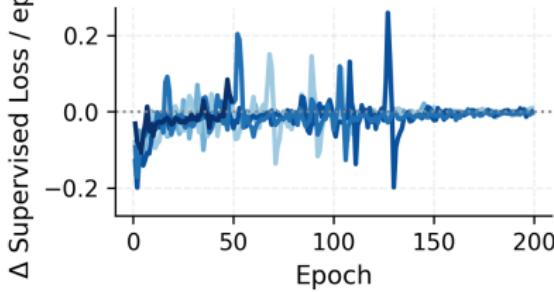
**4 labels**



**25 labels**



**400 labels**

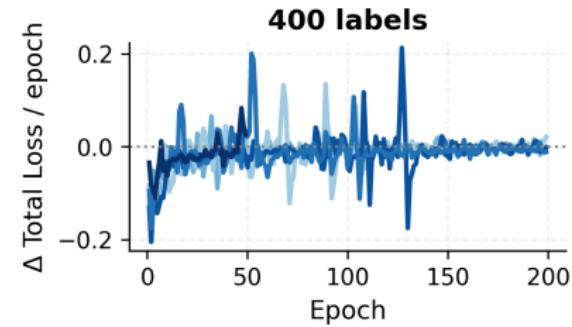
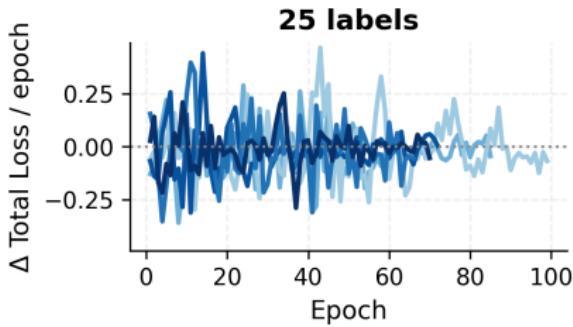
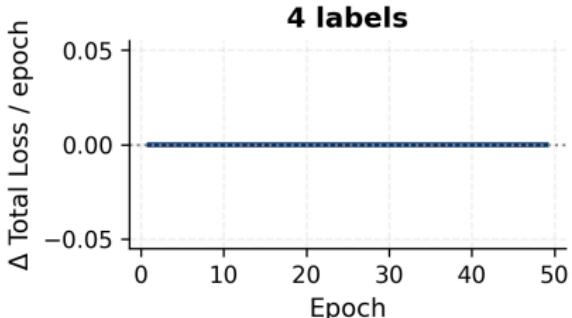
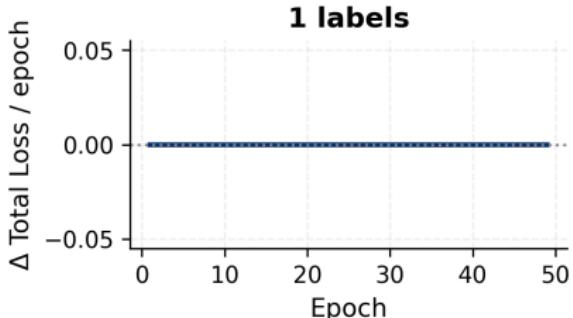


Legend:

- ResNet-18
- ResNet-30
- ResNet-50
- ResNet-101
- ResNet-152

# Change in Total Loss vs Epoch by Label Count

ResNet Family

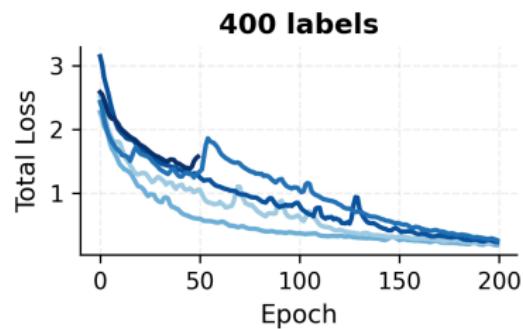
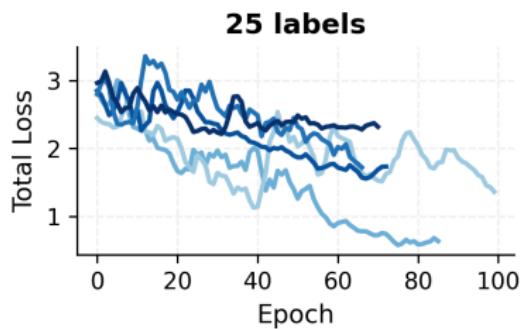
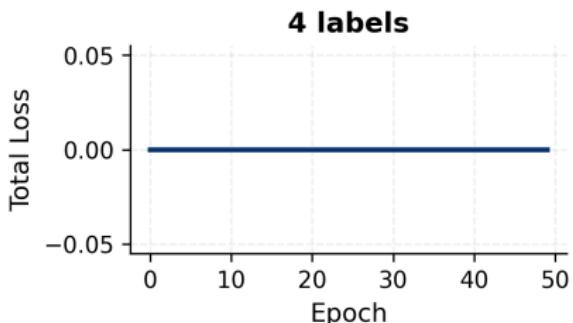
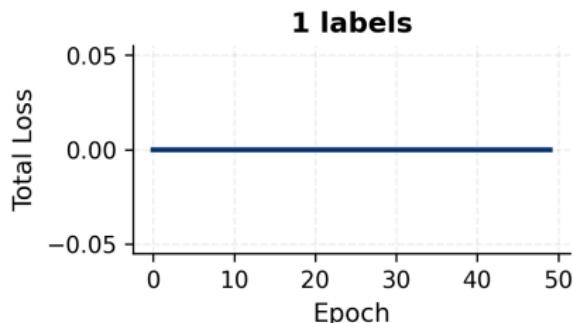


Legend:

- ResNet-18
- ResNet-50
- ResNet-101
- ResNet-152
- ResNet-34

# Total Loss vs Epoch by Label Count

ResNet Family



Legend:

- ResNet-18
- ResNet-34
- ResNet-50
- ResNet-101
- ResNet-152

## Change in Unsupervised Loss vs Epoch by Label Count

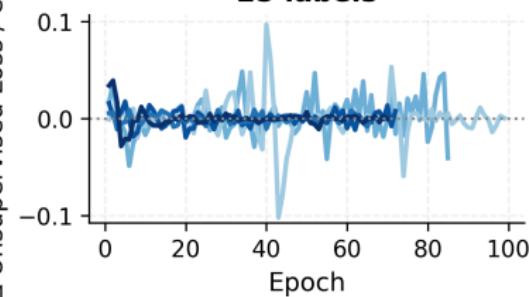
ResNet Family

$\Delta$  Unsupervised Loss / epoch

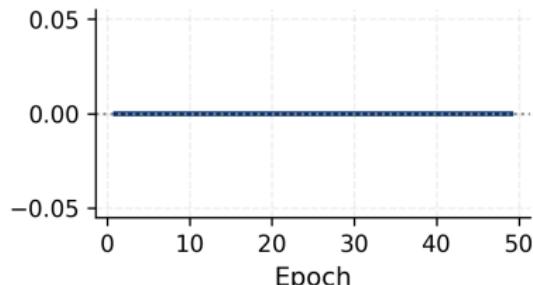
**1 labels**



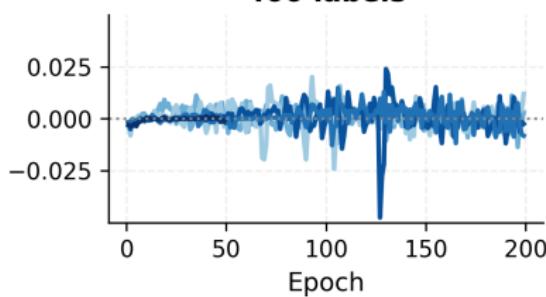
**25 labels**



**4 labels**



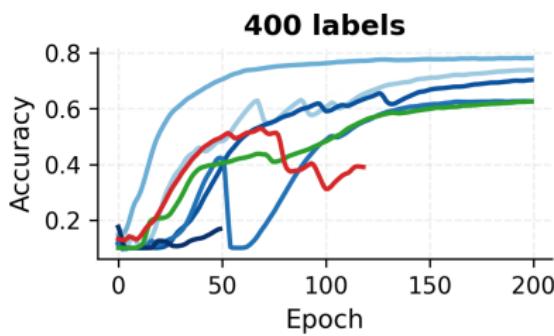
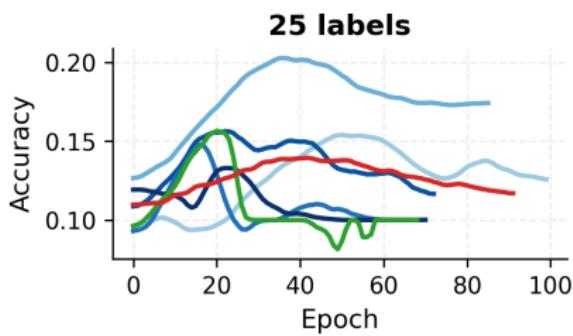
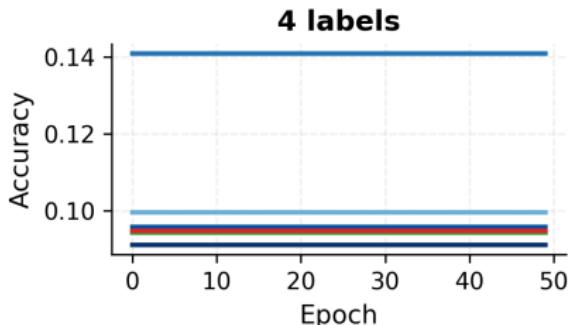
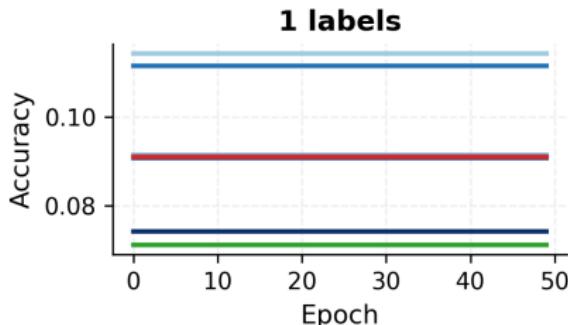
**400 labels**



Legend: ResNet-18 (light blue), ResNet-34 (medium light blue), ResNet-50 (dark blue), ResNet-101 (blue-gray), ResNet-152 (dark blue-gray)

# **ResNet + Outros**

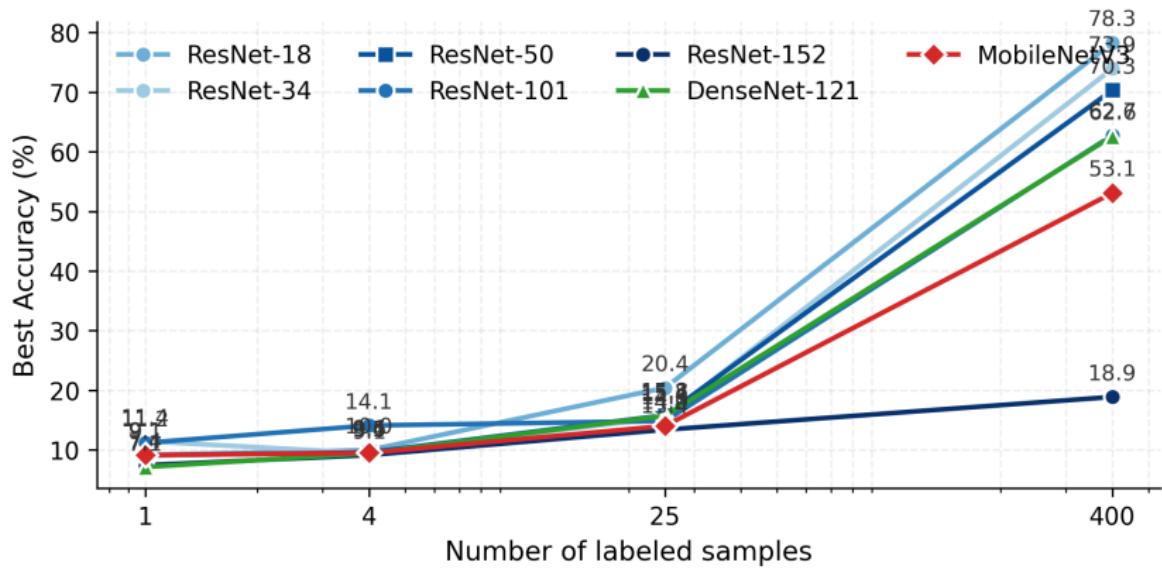
## Accuracy vs Epoch by Label Count

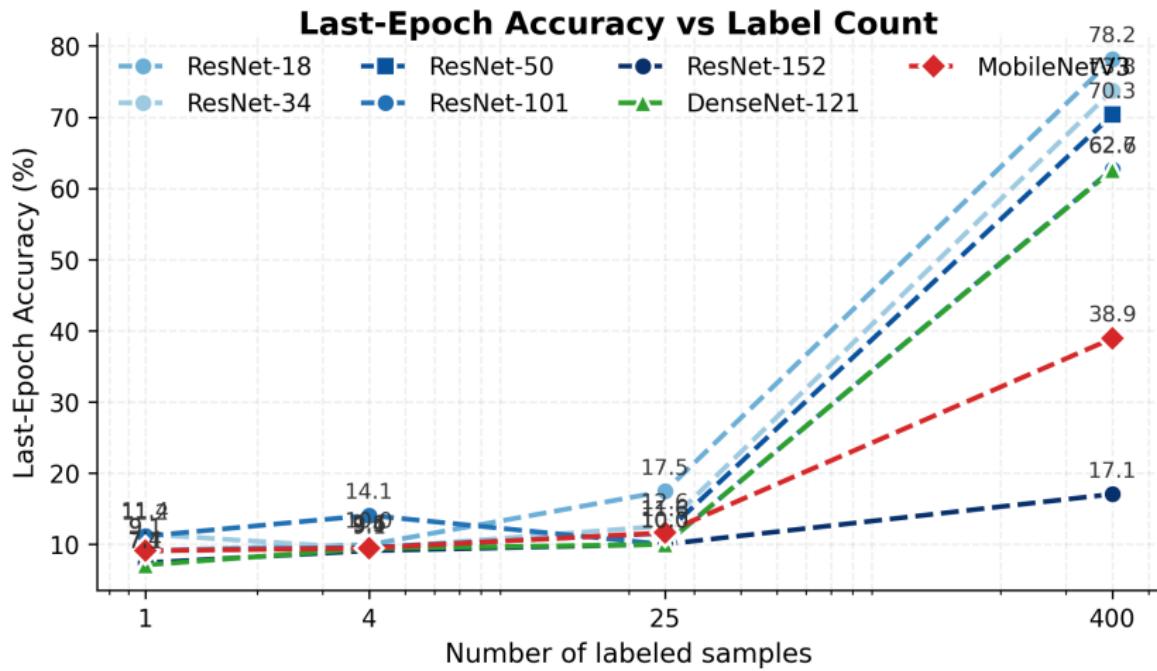


Legend:

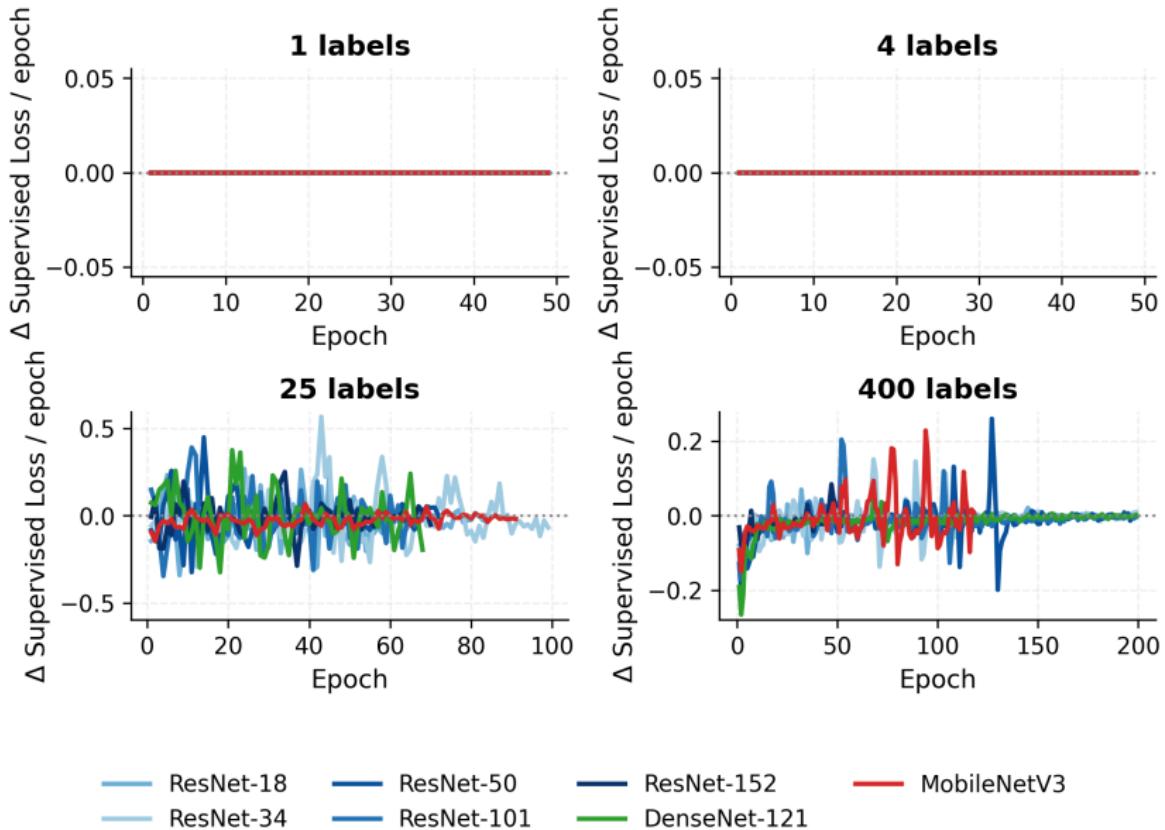
- ResNet-18
- ResNet-34
- ResNet-50
- ResNet-101
- ResNet-152
- DenseNet-121
- MobileNetV3

## Best Accuracy vs Label Count

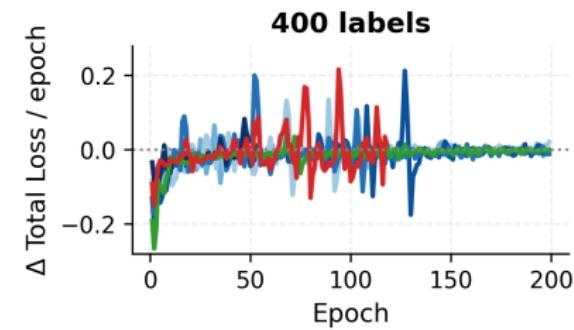
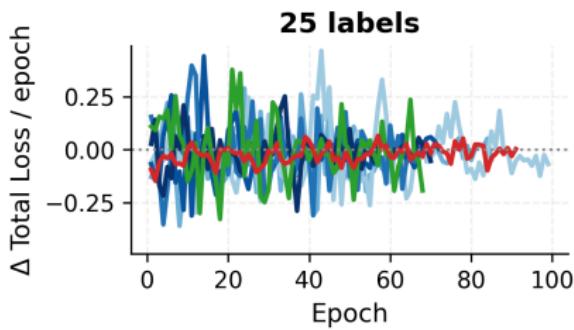
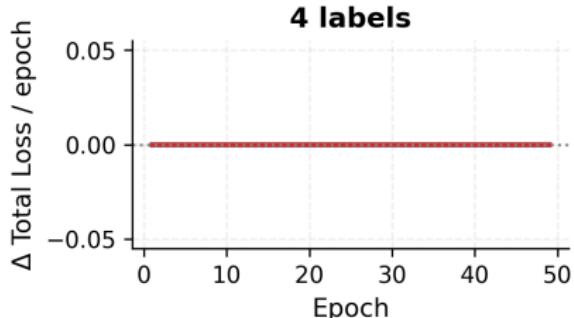
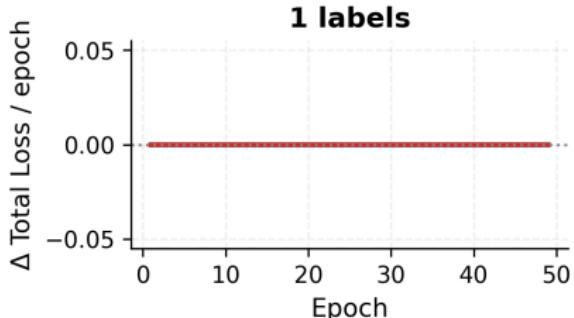




## Change in Supervised Loss vs Epoch by Label Count



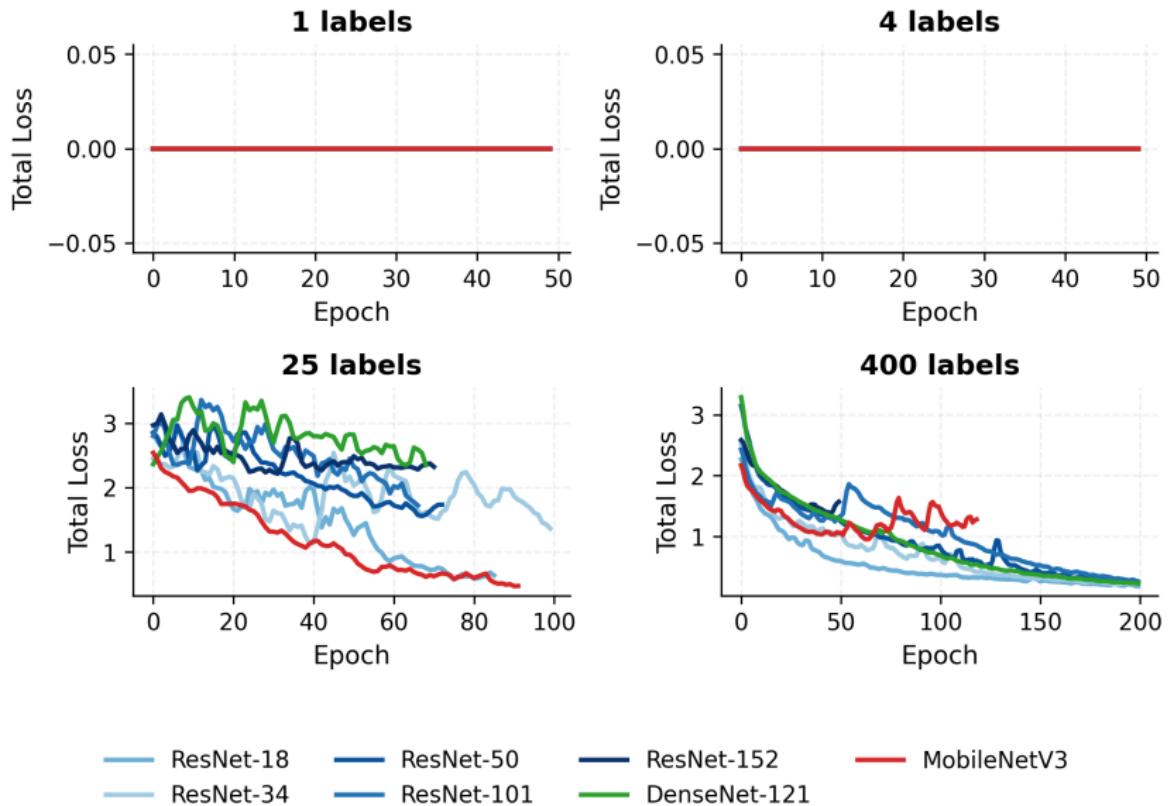
## Change in Total Loss vs Epoch by Label Count



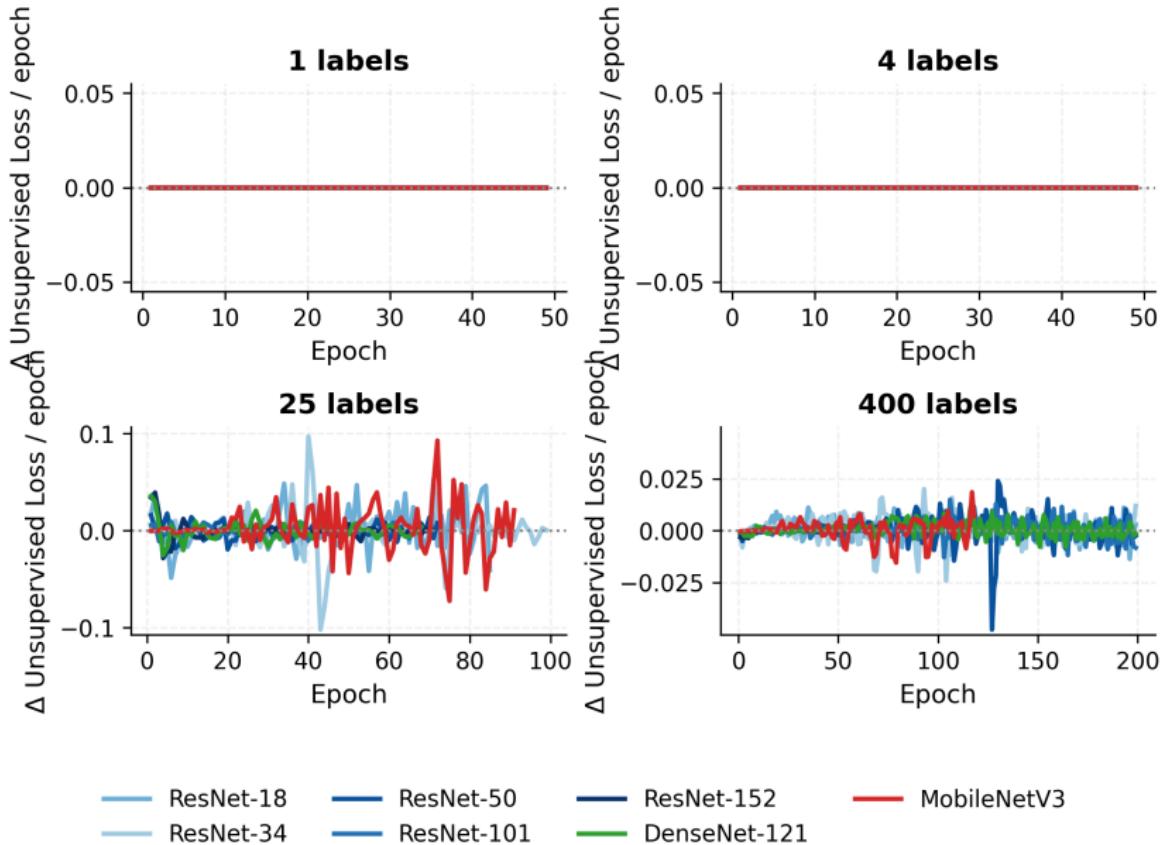
Legend:

ResNet-18	ResNet-50	ResNet-152	MobileNetV3
ResNet-34	ResNet-101	DenseNet-121	

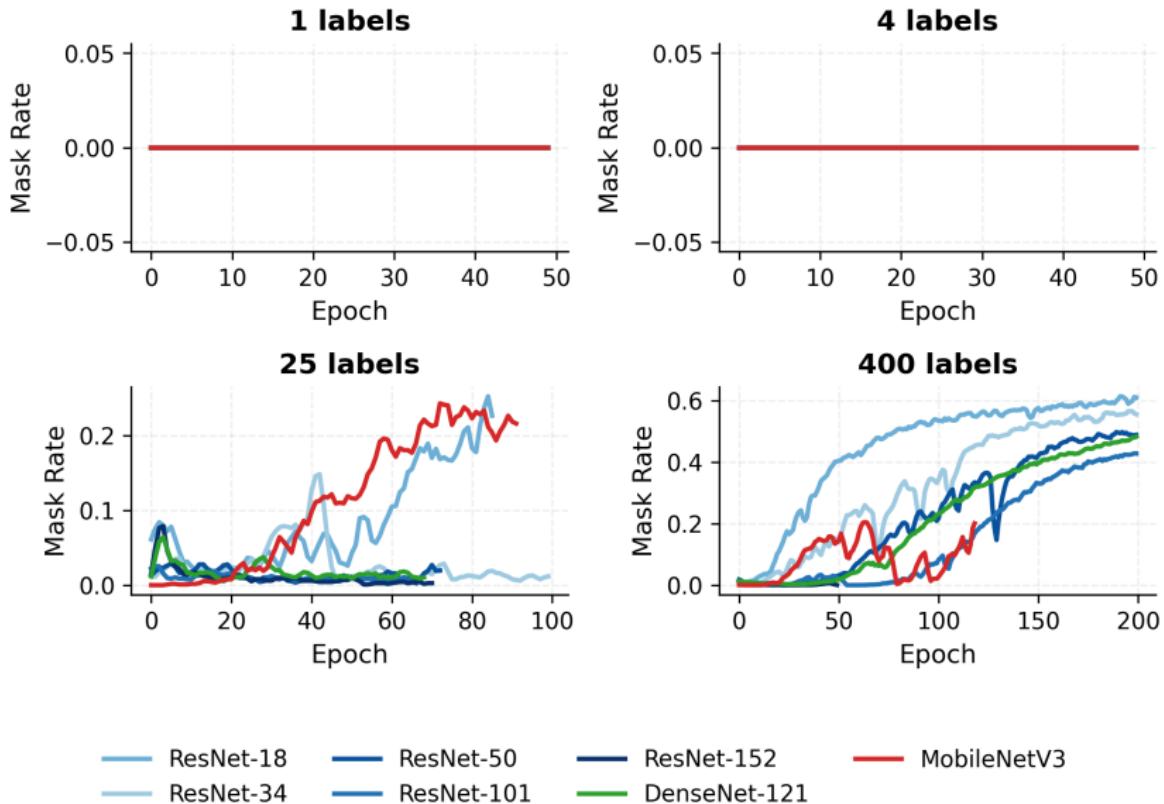
## Total Loss vs Epoch by Label Count



## Change in Unsupervised Loss vs Epoch by Label Count



## Mask Rate vs Epoch by Label Count

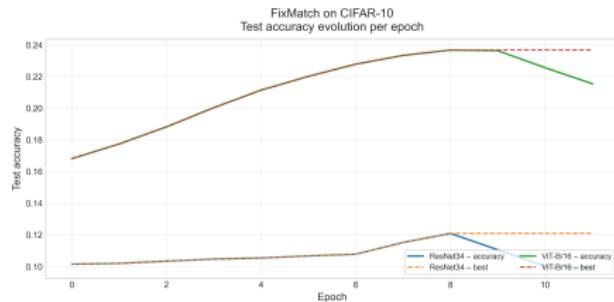


# Por que comparar ViT-B/16 e ResNet34?

- **Contraste de paradigmas:** ResNet34 é o backbone convolucional clássico indicado no paper original do FixMatch; ViT-B/16 representa a família transformer com atenção global.
- **Expectativa a priori:** esperávamos que o ViT atingisse acurácia superior após poucas épocas graças à modelagem de longo alcance, enquanto a ResNet serviria como baseline estável e já conhecido.
- **Importância prática:** comparar ambos mostra o quanto ganhos reais vêm de trocar o backbone, mantendo o mesmo pipeline semi-supervisionado (mesmo threshold, EMA, augmentações).
- **Considerações gerais:** ViTs exigem regularização pesada e maior custo computacional; ResNets convergem mais lentamente, mas demandam menos memória. Os gráficos seguintes quantificam esse trade-off.

# Trajetórias de Acurácia (EMA vs. cru)

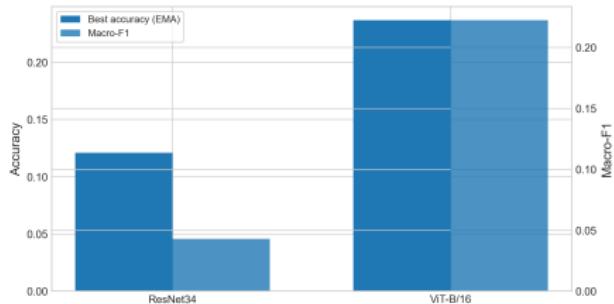
- ViT-B/16 já supera 17% de acurácia de teste na época 0 e cruza 20% até a época 4.
- ResNet34 estaciona próximo de 12% mesmo com o EMA, sem ganhos relevantes após a época 8.



# Melhor Acurácia vs. Macro-F1

- Melhor acurácia EMA:  
ViT-B/16 em 23,7% vs.  
ResNet34 em 12,1% (gap de  
+11,6 pp).
- Macro-F1 repete o cenário:  
22,2% para ViT-B/16 e  
apenas 4,3% para ResNet34.

FixMatch on CIFAR-10  
Comparing ResNet34 vs ViT-B/16



# Gap de Acurácia ao Longo do Treino

- ViT mantém vantagem em todo o treino, atingindo pico de +12 pp perto da época 8.
- Nenhuma época apresenta liderança da ResNet, indicando benefício estrutural do backbone transformer.
- O gap diminui levemente após a época 10, sugerindo convergência sem cruzamento das curvas.

